Deep Q-Learning with Recurrent Neural Networks

Clare Chen
cchen9@stanford.edu

Vincent Ying
vincenthying@stanford.edu

Dillon Laird
dalaird@cs.stanford.edu

Abstract

Deep reinforcement learning models have proven to be successful at learning control policies image inputs. They have, however, struggled with learning policies that require longer term information. Recurrent neural network architectures have been used in tasks dealing with longer term dependencies between data points. We investigate these architectures to overcome the difficulties arising from learning policies with long term dependencies.

1 Introduction

Recent advances in reinforcement learning have led to human-level or greater performance on a wide variety of games (e.g. Atari 2600 Games). However, training these networks can take a long time, and the techniques presented in the state of the art [0] perform poorly on several games that require long term planning.

Deep Q-networks are limited in that they learn a mapping from a single previous state which consist of a small number of game screens. In practice, DQN is trained using an input consisting of the last four game screens. Thus, DQN performs poorly at games that require the agent to remember information more than four screens ago. This is evident from the types of games DQN performs poorly at, near or below human-level [0], in Figure 1.

We explore the concept of a deep recurrent Q-network (DRQN), a combination of a recurrent neural network (RNN) [6] and a deep Q-network (DQN) similar to [5] [1]. The idea being that the RNN will be able to retain information from states further back in time and incorporate that into predicting better Q values and thus performing better on games that require long term planning.

Figure 1: Q*bert, Seaquest, SpaceInvaders and Montesuma's Revenge

In addition to vanilla RNN architectures, we also examine augmented RNN architectures such as attention RNNs. Recent achievements of attention RNNs in translation tasks [2, 3] have shown promise. The advantage of using attention is that it enables DRQN to focus on particular previous states it deems important for predicting the action in the current state. We investigate augmenting DRQN with attention and evaluate its usefulness.

---

[1] code at https://github.com/dillonalaird/deep-rl-tensorflow

2   Related Work

Reinforcement learning covers a variety of areas from playing backgammon [7] to flying RC helicopters [8]. Traditionally, reinforcement learning relied upon iterative algorithms to train agents on smaller state spaces. Later, algorithms such as Q-learning were used with non-linear function approximators to train agents on larger state spaces. These algorithms, however, were more difficult to train and would diverge [9].

Recent advances in reinforcement learning have made it possible to use deep neural networks as non-linear function approximators and train them without running into stability issues. These types of models are known as Q-networks and have been shown to be very successful at playing games such as Atari games where the input is nothing more than the pixels on the screen [0, 1].

We examine extending this framework to include RNNs. RNNs have been used before in Q-Learning [5] but on partially observable Markov decision processes created by flickering the game screen. Our goal is to improve the average score the agent receives.

3   Deep Q-Learning

We examine agents playing Atari games. In our environment, the agents interact with an Atari emulator. At time t they recieve an observation $x_t \in \mathbb{R}^D$, which is a vector of pixels from the current game screen. The agent then takes an action $a_t \in \mathcal{A} = \{1, \ldots, K\}$ and receives a reward r_t which is the change in the game score.

The objective of the agent is to take actions that maximize the future discounted rewards. We can calculate the future rewards with $R_t = \sum_{t'=1}^{T} \gamma^{t'-1} r_{t'}$ where T is the end of the game and γ is our discount factor. One way to achieve this objective is by taking the action corresponding to the maximum action-value function:

$$Q^*(s, a) = \max_{\pi} \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \ldots | s_t = s, a_t = a, \pi],$$

which is the expected value of the future discounted rewards under the policy $\pi = P(a|s)$. This is the probability of taking an action given the agent is in a certain state. Unfortunately, this method can be very difficult to calculate, so we approximate it with another function $Q(s, a; \theta) \approx Q^*(s, a)$. We examine several different types of deep neural network architectures for our function approximator which are called Q-networks.

Training this function can be unstable and will sometimes diverge. The following following loss function is used to help reduce some of these issues [0]:

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right] \tag{1}$$

where experiences are defined as $e_t = (s_t, a_t, r_t, s_{t+1})$ and the buffer in which experiences are stored is $D_t = \{e_1, \ldots, e_t\}$. In the SGD update, we sample mini-batches of experience tuples $(s, a, r, s') \sim U(D)$, which are drawn uniformly at random from the memory buffer. This is called experience replay and is designed to mitigate instability in the training process. In the loss function above, θ_i^- are the network parameters used to compute the target network at iteration i. We only update these target network parameters with the Q-network parameters θ_i every C steps; note that we are keeping a Q-network and a target network.

4   Deep Recurrent Q-Learning

We examined several architectures for the DRQN. One of them is the use of a RNN on top of a DQN, to retain information for longer periods of time. This should help the agent accomplish tasks that may require the agent to remember a particular event that happened several dozens screen back.

We also examined the use of an attention mechanism in RNNs. Attention allows the RNN to focus on particular states it has seen in the past. One can think of it as assigning importance to

the states iterated over by the RNN. We investigate 2 forms of attention, a linear attention that uses a learned vector to assign importances over the previous states and a global attention that assigns importances to previous states based on the current state.

The first architecture is a basic extension of DQN. The architecture of DRQN augments DQN's fully connected layer with a LSTM. We accomplish this by looking at the last L states, $\{s_{t-(L-1)}, \ldots, s_t\}$, and feed these into a convolutional neural network (CNN) to get intermediate outputs $\text{CNN}(s_{t-i}) = x_{t-i}$. These are then fed into a RNN (we use an LSTM for this but it can be any RNN), $\text{RNN}(x_{t-i}, h_{t-i-1}) = h_{t-i}$, and the final output h_t is used to predict the Q value which is now a function of $Q(\{s_{t-(L-1)}, \ldots, s_t\}, a_t)$.


Figure 2: Architecture of the Attention DRQN

The second architecture we used was a version of an attention RNN, we are call linear attention. For the linear attention RNN, we take the L hidden states outputted by the RNN, $\{h_{t-(L-1)}, \ldots, h_t\}$, and calculate an inner product with v_a, $\{v_a^T h_{t-(L-1)}, \ldots, v_a^T h_t\}$. This allows the model to focus more on nearby hidden states or further away states depending on what values of v_a are learned. We then take a softmax over these values, $a_{t-i} = \text{softmax}(v_a^T h_{t-i})$. We use this softmax to take a weighted sum over the hidden states to get a context vector, $c_t = \sum_{i=0}^{L-1} a_{t-i} h_{t-i}$. This context vector is then used to predict the Q value.

The third architecture we used includes a global attention mechanism, similar to the global attention used in [3]. A diagram of this type of attention can be seen in Figure 2. We treat the current state s_t as the "decoder" input and the previous $L - 2$ states as the "encoder" inputs. We compute the following scores, $\{h_{t-(L-1)}^T h_t, \ldots, h_{t-1}^T h_t\}$. We then take a softmax over these values, $a_{t-i} = \text{softmax}(h_{t-i}^T h_t)$. The context vector is computed as a weighted sum over the previous hidden states, $c_t = \sum_{i=1}^{L-1} a_{t-i} h_{t-i}$. Finally, the context vector is used to compute $\tilde{h} = \tanh(W_a[h_t; c_t] + b_a)$, which is then used to predict the Q value. This type of attention allows the model to focus on previous states depending on the current state h_t, as opposed to a fixed vector such as v_a.

Learning sequences of observations creates some difficulty when sampling experiences and using the loss function defined in (1). We propose a simple solution where we sample $e_t \sim U(D)$, take the previous L states, $\{s_{t-(L+1)}, \ldots, s_t\}$, and then zero out states from previous games. For example, if s_{t-i} was the end of the previous game, then we would have states $\{\mathbf{0}, \ldots, \mathbf{0}, s_{t-(i+1)}, \ldots, s_t\}$ and similarly for the next state $\{\mathbf{0}, \ldots, \mathbf{0}, s_{t-(i+2)}, \ldots, s_{t+1}\}$.

5   Experiments

For our experiments, we mainly focused on the game Q*bert. We chose Q*bert because it was a challenging game for DQN which achieved scores only slightly above human-level [0] but it was not so challenging that DQN could not make any progress, such as Montezuma's Revenge [0].

For input, the DRQN takes a specified number of screens per state. The screen images are gray-scaled and resized to 80×80. The first hidden layer convolves 16 filters (19×19) with stride 8 across the input image and applies a rectified linear unit. The second hidden layer convolves 32 filters (8×8) with stride 4, again followed by a rectified linear unit. Convolutional outputs are fed to an LSTM with 256 hidden units per layer. Finally, a fully connected linear layer outputs a Q-value for each possible action. All of the hyperparameters used are listed in appendix A. These are similar to the hyperparameters used in [0].



Figure 3: Graphs for Q*bert Over 5 Million Iterations

Figure 4: Graphs for Seaquest Over 10 Million Iterations

For Figure 3, we trained 4 different models for about 5 million iterations on Q*bert which took about 6 days to complete. All the models use the same convolutional neural network described above for processing the images. DQN uses 4 game screens per state.

The DRQN models in Figure 3 used $L = 16$ and 2 screens per state, which allows it to see the last 32 frames. We wanted to let the model look over enough previous states to make an informed decision, but not so many that RNN suffers from vanishing gradients and training time issues which is why we chose $L = 16$. We found 2 screens per state allows the model to run efficiently without suffering in scores. DRQN is a basic recurrent model with no attention mechanism. DRQN-l is the recurrent model with linear attention and DRQN-g is the recurrent model with global attention.

As you can see from Figure 3, the graphs for the DRQN models are very noisy. The regular DRQN performed the best in terms of average points per life but also had the most difficult time reducing loss. DRQN-l had the noisiest graphs. It performed close to the highest and lowest value at some point on every statistic. DRQN-g performance was disappointing. It had the easiest time minimizing the loss but struggled to gain momentum on average reward per life.

In Figure 4, we trained all 4 models on Seaquest. For our DRQN models, we used $L = 2$ with 2 game screens so they were looking over the same number of game screens as DQN, 4. The average scores all converge to the same value, which is what we would expect since each agent is given the same number of screens with which to predict the Q value.

In Figure 5, we can see the best scores received from each of the algorithms playing 100 games of Q*bert. This table reflects where the algorithms ended up in the average rewards per life graph in Figure 3. DRQN performed the best, followed by DQN, DRQN-l, and finally DRQN-g. Unfortunately, we were not able to reproduce the score achieved by [1] on Q*bert.

	Q*bert	Seaquest
DQN	700	360
DRQN	**850**	360
DRQN-l	700	**620**
DRQN-g	550	380

Figure 5: Best Scores for the Atari Games

From these figures, it is evident that adding a recurrent layer helped the agent, adding an attention mechanism only hindered the agent. We hypothesize several reasons that would explain the unexpected results. For one, the network that must be trained is larger, with more parameters to tune. It is apparent from the graphs that the attention DRQN did not converge. We postulate that training for a longer duration would yield better results. Another reason is that attention might not be necessary for playing Q*bert. Since the agent receives the full game state on each screen, it does not need to focus on game screens too far back in the past.

Conclusion

We investigated the utilization of an RNN on top of a DQN (DRQN) and also tested several different types of attention mechanisms to augment the RNN. As expected, we confirmed that a basic DRQN is capable of achieving better scores than a DQN on a game that is difficult for a DQN to learn. We also discovered that adding attention can hinder the performance on certain games as well.

Going forward, it would be desirable to train the algorithms for even longer to see if attention eventually does perform better. It can be seen that the graphs have not converged yet and are still noisy. Also, we theorize that running the agent on games where the game state is not fully revealed in the screen would better utilize attention to achieve higher scores. Think of a game where in one screen you see a person and in the next screen they run behind a wall. The agent should attend to the screen where the person was last visible to gain information about where their current position might possibly be in the current timestep.

Acknowledgements

Appendix A:

List of Hyperparameters		
minibatch size	32	number of experiences for SGD update
replay memory buffer size	1000000	number of experiences in memory buffer
agent history length	4-32	number of most recent frames experienced input to the Q network
target network update frequency	10000	number of parameter updates after which the target network updates
discount factor	0.99	Discount factor γ used in the Q-learning update
action repeat	4	Repeat each action selected by the agent this many times
update frequency	4	number of actions by agent between successive SGD updates.
initial exploration	1	Initial value of ϵ in ϵ-greedy exploration
final exploration	0.01	Final value of ϵ in ϵ-greedy exploration
final exploration frame	100000	The number of frames over which the initial value of ϵ reaches final value
no-op max	30	Maximum "do nothing" actions by agent at the start of episode.
learning rate of training	0.00025	learning rate used by RMSProp
decay of RMSProp optimizer	0.99	decay rate used by RMSProp
β of RMSProp optimizer	0.01	β rate of RMSProp

References

[0] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level Control through Deep Reinforcement Learning. *Nature*, 518(7540):529-522, 2015.

[1] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra and Martin Riedmiller. Playing Atari with Deep Reinforcement Learning. *arXiv preprint arXiv:1312.5602*, 2013.

[2] Dzmitry Bahdanau, Kyunghyun Cho and Yoshua Bengio. Neural Machine Translation by Jointly Learning to Align and Translate. *arXiv preprint arXiv:1409.0473*, 2014.

[3] Minh-Thang Luong, Hieu Pham and Chisopher D. Manning. Effective Approaches to Attention-based Neural Machine Translation. *arXiv preprint arXiv:1508.04025*, 2015.

[4] Ivan Sorokin, Alexey Seleznev, Mikhail Pavlov, Aleksandr Fedorov and Anastasiia Ignateva. Deep Attention Recurrent Q-Network. *arXiv preprint arXiv:1512.01693*, 2015.

[5] Matthew Hausknecht and Peter Stone. Deep Recurrent Q-Learning for Partially Observable MDPs. *arXiv preprint arXiv:1507.06527*, 2015.

[6] Sepp Hochreiter and Jurgen Schmidhuber. Long Short-Term Memory. *Neural Computation*, 9(8):1735-1780, 1997.

[7] Gerald Tesauro. Temporal difference learning and td-gammon. *Communications of the ACM*, 38(3):58-68, 1995.

[8] Andrew Y. Ng, H. Jin Kim, Michael I. Jordan and Shankar Sastry. Autonomous Helicopter Flight via Reinforcement Learning. *Advances in Neural Information Processing Systems 16*, 2003.

[9] John N. Tsitsiklis and Benjamin Van Roy. An analysis of temporal-difference learning with function approximation. *Automatic Control, IEEE Transactions on*, 42(5):647-690, 1997.